

DD-Mamba: Dual-Domain State-Space Modeling with Zero-Initialized Linear Anchors for Long-Term Time Series Forecasting

Anonymous Submission

ADMA 2026 double-blind review

Abstract. Long-term multivariate time series forecasting must reconcile two complementary views of a signal: local, trend-driven dynamics that live in the time domain, and global periodic structure that is most compact in the frequency domain. Existing state-space (Mamba) forecasters model one view well but not both, and deep forecasters in general discard the simple linear maps that remain near state-of-the-art on several benchmarks. We present DD-Mamba, a dual-domain architecture in which a time branch (trend-seasonal decomposition with a selective state-space encoder) and a frequency branch (learnable complex spectral filtering) each produce their own forecast, fused by a learned convex gate so that no prediction can bypass either domain. Two design principles drive its accuracy: (i) *linear anchoring* — each branch carries an explicit linear pathway of its domain (DLinear-style projections in time, a FITS-style spectral map in frequency) while all nonlinear heads are zero-initialized, so the end-to-end map at initialization is purely linear and nonlinear corrections are learned on top; and (ii) *capacity placement* — a bidirectional Mamba over the *variate* dimension carries the model on high-channel data and is disabled on small, few-channel data where it only adds overfitting capacity. On nine benchmarks under the look-back-96 protocol, DD-Mamba attains the lowest average MSE on seven and improves over the strongest state-space baseline S-Mamba on seven (up to 3.5% on ETTh1) at a fraction of the parameters on small datasets. Ablations support each design choice and surface a cautionary finding — instance normalization *helps* traffic yet *hurts* solar — so normalization must be validated per dataset.

Keywords: Time series forecasting · State-space models · Mamba · Frequency domain · Multivariate analysis

1 Introduction

Long-term forecasting of multivariate time series underpins applications from power-grid load planning to traffic management and weather-dependent operations. The field has cycled through architectural paradigms at a remarkable pace: Transformer variants with efficient attention [1–3], patch-based attention [4], inverted variate-token attention [5], and, most recently, selective state-space

models (SSMs) built on Mamba [11, 12], which offer linear-time sequence modeling with input-dependent dynamics.

Two empirical facts complicate this progress. First, *simple linear maps remain stubbornly competitive*: a per-component linear projection from the look-back window to the horizon (DLinear [6]) and a tiny complex-linear interpolation of the low-passed spectrum (FITS [10]) match or beat far larger models on several benchmarks, especially the ETT family. Deep architectures that ignore this fact spend their early training re-discovering a linear map, and on small datasets often never do. Second, *the useful inductive bias differs by domain*: local trends and regime changes are naturally expressed in the time domain, while multi-scale seasonality is compact in the frequency domain. Models committed to a single view — attention or SSMs over time steps only, or spectral mixing only [9] — systematically underuse the other.

We address both issues with DD-Mamba, a dual-domain forecaster organized around three ideas.

Dual branches that cannot be bypassed. A time branch and a frequency branch each output a complete forecast; the model’s prediction is a learned *convex combination* of the two, with the gate conditioned on both branches’ features. Because the output is constrained elementwise between the branch forecasts, no shortcut path around either domain exists — a property we verify by construction and exploit for per-sample attribution of the two domains.

Linear anchors with zero-initialized corrections. Each branch carries an explicit *linear* pathway that parameterizes the strongest linear forecaster family of its domain: seasonal/trend window-to-horizon projections [6] in the time branch, and a FITS-style complex-linear spectral map [10] in the frequency branch. All nonlinear heads (and the spectral map) are zero-initialized and the fusion gate opens as a constant, so the end-to-end function at initialization is *purely linear* — a randomly initialized DLinear scaled by the gate — and gradient descent learns the spectral anchor and the nonlinear corrections jointly, only where the data supports them (Sect. 3.2). This choice stabilizes small-data training and removes the mean-reversion failure mode of pooled deep forecasters.

Mamba where it pays. A causal Mamba encoder models the time axis; a *bidirectional* Mamba (frequency bins and channels have no arrow of time) serves two optional roles: spectral encoding, and cross-channel mixing over the variate dimension in the spirit of S-Mamba [12]. Crucially, capacity is *placed* per dataset: the variate mixer is the dominant component on high-channel data (electricity, traffic, solar) and is disabled on few-channel small data (ETT, exchange) where ablations show it only adds overfitting capacity.

Our contributions are:

1. A dual-domain forecasting architecture whose output is a gated convex combination of per-branch forecasts, guaranteeing that both the time and the frequency view contribute to every prediction (Sect. 3).
2. A linear-anchoring scheme — domain-appropriate linear pathways with zero-initialized nonlinear corrections — under which the end-to-end map at ini-

tialization is purely linear, yielding large gains on small benchmarks (Sect. 3.2, ablated in Sect. 4.3).

3. A capacity-placement recipe for selective SSMS: causal Mamba over time where temporal microstructure matters, bidirectional Mamba over variates where cross-channel structure dominates, and neither where data is scarce (Sect. 3.4).
4. The lowest average MSE on seven of nine standard benchmarks under the unified look-back-96 protocol, plus a released ablation harness and a case-study ablation that isolates every component — including negative and dataset-dependent findings (instance normalization helps traffic by +0.10 MSE yet hurts solar by 0.016) that we believe are as actionable as the positive ones (Sect. 4).

2 Related Work

Transformer forecasters. Informer [1], Autoformer [2] and FEDformer [3] adapt attention to long sequences via sparsity, decomposition, and frequency-domain mixing respectively. PatchTST [4] treats sub-series patches as tokens with channel independence; iTransformer [5] inverts the axes and attends over *variates*, establishing that cross-channel modeling, not temporal attention, drives accuracy on high-dimensional benchmarks. DD-Mamba adopts the variate-mixing insight but implements it with linear-time bidirectional SSMS and, unlike these works, couples it to an explicit frequency branch and linear anchors.

Linear and frequency-domain models. DLinear/NLinear [6] showed that decomposition plus a linear map rivals Transformers, particularly on ETT; RLinear [7] added reversible instance normalization [8]. FreTS [9] learns MLPs in the frequency domain, and FITS [10] forecasts with a single complex linear interpolation of the low-passed spectrum at negligible parameter cost. DD-Mamba does not compete with these models — it *contains* them: each branch’s backbone is the linear model of its domain, and the deep machinery is trained as a correction on top.

State-space models. Mamba [11] introduced selective SSMS with input-dependent dynamics and linear-time scans. In forecasting, S-Mamba [12] applies bidirectional Mamba over variate tokens with a linear temporal embedding, and further works explore frequency-domain SSMS [13]. DD-Mamba unifies these threads: causal Mamba over time, bidirectional Mamba over frequency bins and over variates, with the placement of each governed by dataset characteristics rather than applied uniformly.

3 Method

3.1 Problem Formulation and Normalization

Given a look-back window $\mathbf{X} \in \mathbb{R}^{L \times C}$ of L steps and C variates, the task is to predict the next H steps $\mathbf{Y} \in \mathbb{R}^{H \times C}$. Following RevIN [8], we standardize

each window per variate, $\tilde{\mathbf{X}} = (\mathbf{X} - \boldsymbol{\mu})/\boldsymbol{\sigma}$, and de-normalize the final forecast. Instance normalization is treated as a per-dataset switch rather than a universal default; Sect. 4.3 shows why.

3.2 Time Branch: Decomposition, Linear Backbone, Selective Scan

The time branch decomposes $\tilde{\mathbf{X}}$ with a moving average into trend $\mathbf{X}_t$ and seasonal $\mathbf{X}_s$ components [2, 6] and forecasts with two cooperating paths.

Linear backbone. Per-component linear maps project the window directly to the horizon:

$$\mathbf{Y}_{\text{lin}} = \mathbf{W}_s \mathbf{X}_s^\top + \mathbf{W}_t \mathbf{X}_t^\top, \quad \mathbf{W}_s, \mathbf{W}_t \in \mathbb{R}^{H \times L}, \quad (1)$$

i.e., exactly DLinear [6], shared across variates.

Selective-scan correction. The pair $[\mathbf{X}_s; \mathbf{X}_t]$ is embedded per time step and processed channel-independently by a stack of Mamba layers [11]. A Mamba layer discretizes a state-space system with input-dependent parameters $(\boldsymbol{\Delta}, \mathbf{B}, \mathbf{C})$:

$$\mathbf{h}_k = e^{\boldsymbol{\Delta}_k \mathbf{A}} \mathbf{h}_{k-1} + \boldsymbol{\Delta}_k \mathbf{B}_k u_k, \quad y_k = \mathbf{C}_k^\top \mathbf{h}_k + D u_k, \quad (2)$$

scanned causally over the L steps. The last state, summed with a whole-window linear embedding $\mathbf{W}_e \tilde{\mathbf{x}}^{(e)}$ (one per variate, as in [12]), forms the branch feature $\mathbf{f}_{\text{time}} \in \mathbb{R}^{C \times d}$. A head ϕ_{time} maps features to a horizon correction:

$$\mathbf{Y}_{\text{time}} = \mathbf{Y}_{\text{lin}} + \phi_{\text{time}}(\mathbf{f}_{\text{time}}), \quad \phi_{\text{time}} \equiv \mathbf{0} \text{ at init.} \quad (3)$$

Because ϕ_{time} is zero-initialized, the branch *is* DLinear at initialization; the SSM learns only what the linear map cannot express. We compute the recurrence of Eq. (2) in fp32 even under mixed-precision training: the exponential and the L -step state accumulation overflow in half precision, which silently destroys training (observed as NaN validation losses that exhaust early-stopping patience).

3.3 Frequency Branch: Spectral Filtering and a FITS Anchor

The frequency branch applies the real FFT along time, $\mathbf{Z} = \text{rFFT}(\tilde{\mathbf{X}}) \in \mathbb{C}^{C \times F}$, $F = \lfloor L/2 \rfloor + 1$, optionally low-passed by zeroing the top fraction ρ of bins.

Spectral encoder. A learnable complex linear filter $\mathbf{W} = \mathbf{W}_r + i\mathbf{W}_i$ densely mixes bins, $\mathbf{Z}' = \mathbf{W}\mathbf{Z}$, realized with two real matrices; a bidirectional Mamba over the bin sequence is available as an input-dependent alternative. The (real, imaginary) parts are projected to the branch feature $\mathbf{f}_{\text{freq}} \in \mathbb{R}^{C \times d}$ and a head ϕ_{freq} produces the spectral forecast.

FITS anchor. Mirroring the time branch, a complex linear map $\mathbf{U} \in \mathbb{C}^{F' \times F}$ with $F' = \lfloor (L+H)/2 \rfloor + 1$ maps the (low-passed) window spectrum to that of a length- $(L+H)$ series, whose inverse transform’s last H samples are the branch’s linear spectral forecast:

$$\mathbf{Y}_{\text{freq}} = \underbrace{\text{irFFT}_{L+H}(\mathbf{U}\mathbf{Z})[L:]}_{\text{linear-in-frequency}} + \phi_{\text{freq}}(\mathbf{f}_{\text{freq}}), \quad \mathbf{U}, \phi_{\text{freq}} \equiv \mathbf{0} \text{ at init.} \quad (4)$$

This pathway *parameterizes* the FITS [10] family — any linear-in-frequency forecaster of the low-passed spectrum is reachable — but, unlike FITS itself, $\mathbf{U}$ is *zero-initialized*, so the frequency branch is silent at initialization rather than a functioning spectral predictor. This is deliberate: the frequency grids of a length- L and a length- $(L+H)$ window do not align (harmonic k of the input maps to the non-integer index $k(L+H)/L$ of the output), so no canonical identity or interpolation initialization exists; a randomly initialized $\mathbf{U}$ would inject noise into the fused forecast during the highest-learning-rate epochs. Zero initialization instead lets the gate lean on the time anchor until $\mathbf{U}$ is learned jointly with the rest of the model. Consequently, the precise initialization claim is: the time branch coincides with a (randomly initialized) DLinear, the frequency branch outputs zero, and the fusion gate is the constant g_0 (Sect. 3.5), so the *end-to-end* map is the purely linear forecaster $g_0 \cdot \mathbf{Y}_{\text{lin}}$. Both branches retain explicit linear pathways of their domain throughout training.

3.4 Cross-Channel Mixing over Variates

Channel independence is a strong baseline [4, 6] but leaves accuracy on the table when variates are strongly coupled [5, 12]. Inside each branch, before its head, we optionally run M layers of *bidirectional* Mamba over the variate dimension — treating the C per-channel features as a sequence — followed by a position-wise feed-forward block. Bidirectionality is essential: channel order carries no causal direction, so each layer runs forward and reversed scans and sums them in one residual stream. The mixer is a per-dataset choice: $M=2$ at width 512 on electricity/traffic (where it is the single most load-bearing component; Sect. 4.3), $M=2$ at width 256 on weather, $M=1$ on solar, and $M=0$ on the 7–8-channel ETT and exchange datasets, where it triples parameters without adding signal.

3.5 Forecast-Level Convex Fusion

Each branch emits a forecast and a feature. The model output is a gated convex combination

$$\mathbf{Y} = g \odot \mathbf{Y}_{\text{time}} + (1 - g) \odot \mathbf{Y}_{\text{freq}}, \quad g = \sigma(\psi([\mathbf{f}_{\text{time}}; \mathbf{f}_{\text{freq}}])) \in (0, 1)^{C \times 1}, \quad (5)$$

followed by RevIN de-normalization. Since $\mathbf{Y}$ lies elementwise between the branch forecasts, no path can bypass the dual-domain structure; the gate also provides free per-sample attribution of the two domains. ψ is zero-initialized with bias 2.2, so training starts at $g \approx 0.9$: the time branch is an exact linear model at initialization while the frequency head is silent, and an even 0.5/0.5 mix would waste the highest-learning-rate epochs compensating for noise.

3.6 Training

We minimize MSE with AdamW under a halving schedule ($\eta_e = \eta_0 \cdot 2^{-(e-1)}$), which concentrates learning in the first epochs and suppresses the late-epoch overfitting that cosine schedules invite on small benchmarks. All runs use 10 epochs; early stopping monitors validation loss with the best checkpoint restored.

Table 1. Benchmark datasets and where DD-Mamba places its capacity.

Dataset	Variates	Steps	Granularity	Variate mixer	Params
ETTh1/ETTh2	7	17,420	1 hour	off	0.25M
ETTm1/ETTm2	7	69,680	15 min	off	0.24–0.36M
Weather	21	52,696	10 min	$M=2$, $d=256$	5.8M
Solar-Energy	137	52,560	10 min	$M=1$, $d=128$	0.96M
Electricity	321	26,304	1 hour	$M=2$, $d=512$	19.7M
Traffic	862	17,544	1 hour	$M=2$, $d=512$	19.7M
Exchange	8	7,588	1 day	off	0.09M

4 Experiments

4.1 Setup

Datasets and protocol. We evaluate on nine standard benchmarks (Table 1) under the unified protocol of [5, 12]: look-back $L=96$ for all datasets and horizons $H \in \{96, 192, 336, 720\}$; ETT uses the canonical 12/4/4-month border split (the file tail beyond month 20 is unused), the others chronological 0.7/0.1/0.2 splits; variates are standardized by training-set statistics; metrics are MSE and MAE on the standardized test set.

Baselines. We compare against the strongest published results under the same protocol: the state-space forecaster S-Mamba [12], iTransformer [5], PatchTST [4], DLinear [6], and TimesNet [14]. Baseline numbers are quoted from [12] rather than reproduced: the evaluation protocol (look-back, splits, standardization, metrics) is identical, but training pipelines, seeds, and implementations differ across works, so sub-noise-floor gaps to quoted numbers should be read with caution. We flag every such case in the text, and re-running all baselines within one unified training framework is planned for the camera-ready appendix.

Implementation and Reproducibility. All experiments use PyTorch 2.11 (CUDA 13.0) on a single NVIDIA RTX 5090 (32 GB). Optimization is AdamW with gradient clipping at 5.0, 10 epochs with the learning rate halved after every epoch, early stopping on validation loss with the best checkpoint restored, and mixed precision with the selective scan forced to fp32 (Sect. 3.2). The decomposition kernel is 25; Mamba blocks use expansion 2 and depthwise convolution width 4 throughout. Table 2 lists all per-dataset settings; complete configuration files and the ablation harness ship with the code. Main results (Tables 3 and 4) are single runs with seed 2024; from repeated runs we estimate a run-to-run noise floor of ≈ 0.005 MSE and treat all differences below it as ties throughout the paper. Reporting multi-seed mean $\pm$ std for every entry is deferred to the camera-ready version; the released harness supports it directly (`-seeds`).

Table 2. Per-dataset hyperparameters. d : model width; time enc.: time-branch encoder (Mamba layers in parentheses); M : variate-mixer layers; N : SSM state size; ρ : spectral low-pass fraction (Sect. 3.3); FITS: spectral linear anchor; IN: instance normalization (RevIN); do.: dropout; B: batch size; wd: weight decay.

Dataset	d	time enc.	M	N	ρ	FITS	IN	lr	B	do./wd
ETTh1	128	Mamba (1)	0	16	0.4	yes	yes	$5 \cdot 10^{-4}$	32	$0.2/10^{-4}$
ETTh2	128	Mamba (1)	0	16	0.3	yes	yes	$5 \cdot 10^{-4}$	32	$0.3/5 \cdot 10^{-4}$
ETTm1	128	Mamba (2)	0	16	0	no	yes	10^{-4}	32	$0.1/10^{-4}$
ETTm2	128	Mamba (1)	0	16	0.2	no	yes	10^{-4}	32	$0.2/2 \cdot 10^{-4}$
Weather	256	Mamba (2)	2	16	0	yes	yes	$2 \cdot 10^{-4}$	32	$0.1/10^{-4}$
Solar	128	Mamba (2)	1	16	0	no	no	$5 \cdot 10^{-4}$	16	$0.1/10^{-4}$
Electricity	512	MLP	2	32	0	no	yes	$5 \cdot 10^{-4}$	16	$0.1/10^{-4}$
Traffic	512	MLP	2	32	0	no	yes	10^{-3}	16	$0.1/10^{-4}$
Exchange	64	Mamba (1)	0	8	0.5	no	yes	10^{-4}	32	$0.3/5 \cdot 10^{-4}$

Table 3. Long-term forecasting results averaged over horizons $H \in \{96, 192, 336, 720\}$ with look-back $L=96$ (MSE/MAE; lower is better). Baselines quoted from [12]. Best in **bold**, second best underlined; exact ties are marked equally.

Dataset	DD-Mamba (ours)	S-Mamba	iTransformer	PatchTST	DLinear	TimesNet
ETTh1	0.439/0.434	0.455/0.450	<u>0.454/0.447</u>	0.469/0.454	0.456/0.452	0.458/0.450
ETTh2	0.375/0.402	<u>0.381/0.405</u>	0.383/0.407	0.387/0.407	0.559/0.515	0.414/0.427
ETTm1	0.386/0.398	0.398/0.405	0.407/0.410	<u>0.387/0.400</u>	0.403/0.407	0.400/0.406
ETTm2	0.280/0.325	0.288/0.332	0.288/0.332	<u>0.281/0.326</u>	0.350/0.401	0.291/0.333
Weather	0.248/0.275	<u>0.251/0.276</u>	0.258/0.278	0.259/0.281	0.265/0.317	0.259/0.287
Electricity	0.169/0.265	<u>0.170/0.265</u>	0.178/0.270	0.205/0.290	0.212/0.300	0.192/0.295
Traffic	0.439/ <u>0.282</u>	0.414/0.276	<u>0.428/0.282</u>	0.481/0.304	0.625/0.383	0.620/0.336
Exchange	0.363/0.405	0.367/0.408	<u>0.360/0.403</u>	0.367/0.404	0.354/0.414	0.416/0.443
Solar	0.228/0.267	0.240/0.273	<u>0.233/0.262</u>	0.270/0.307	0.330/0.401	0.301/0.319

4.2 Main Results

Table 3 summarizes averages over the four horizons; Table 4 reports DD-Mamba per horizon. DD-Mamba attains the lowest average MSE on seven of the nine datasets (ETTh1, ETTh2, ETTm1, ETTm2, Weather, Electricity, Solar) and ranks third on the remaining two. The margins over the runner-up vary widely and we read them accordingly: on ETTh1 and Solar they are substantial (-0.015 and -0.005 MSE), on ETTm1/ETTm2 the advantage over PatchTST is 0.001 — within run-to-run noise (Sect. 4.1), so we regard these as on par — and on Electricity the MAE ties S-Mamba exactly. On Exchange, a near random-walk dataset where all methods cluster within 0.013 MSE, plain DLinear remains best and DD-Mamba is third; on Traffic, S-Mamba and iTransformer lead and DD-Mamba trails by $+0.025$ MSE, which we attribute to the weekly cycle exceeding the protocol look-back (Sect. 4.3). Against S-Mamba specifically, the strongest published state-space forecaster, DD-Mamba has lower average MSE on seven datasets, is on par on Electricity (-0.001), and is behind only on Traffic. The largest margins appear on the small ETT datasets (e.g. ETTh1 0.439 vs. 0.455, -3.5%), where the linear anchors and deliberate capacity restraint matter most; there DD-Mamba uses roughly 0.25M parameters, two orders of magnitude fewer than variate-attention models configured for those datasets.

Table 4. DD-Mamba per-horizon results (MSE/MAE), look-back $L=96$.

Dataset	$H=96$	$H=192$	$H=336$	$H=720$	Avg
ETTh1	0.381/0.397	0.429/0.423	0.472/0.446	0.472/0.468	0.439/0.434
ETTh2	0.292/0.345	0.368/0.392	0.415/0.427	0.423/0.442	0.375/0.402
ETTm1	0.321/0.359	0.366/0.385	0.396/0.405	0.461/0.443	0.386/0.398
ETTm2	0.177/0.260	0.241/0.302	0.302/0.341	0.398/0.397	0.280/0.325
Weather	0.162/0.207	0.212/0.254	0.267/0.293	0.349/0.347	0.248/0.275
Electricity	0.140/0.235	0.157/0.252	0.175/0.271	0.203/0.300	0.169/0.265
Traffic	0.404/0.267	0.423/0.275	0.443/0.283	0.484/0.302	0.439/0.282
Exchange	0.086/0.206	0.179/0.301	0.329/0.415	0.858/0.699	0.363/0.405
Solar	0.197/0.242	0.228/0.265	0.244/0.280	0.243/0.281	0.228/0.267

Table 5. Component ablation on Solar-Energy ($H=96$, single seed; noise floor ≈ 0.005 MSE). Each variant flips exactly one switch of the full configuration.

Variant	MSE	Δ MSE	Component isolated
full	0.197	—	reference
w/o channel mixer	0.219	+0.022	cross-variate Mamba mixing
w/o linear backbone	0.209	+0.012	DLinear anchor (time)
freq branch only	0.209	+0.012	time branch removed
sum fusion	0.202	+0.006	learned gate $\rightarrow$ average
time branch only	0.198	+0.002	frequency branch removed
Mamba $\rightarrow$ MLP over time	0.197	+0.001	selective scan (time)
w/o instance norm	0.181	−0.016	RevIN

4.3 Component Ablations: A Case Study

We release an ablation harness in which every variant flips exactly one switch of a dataset’s full configuration, holding the split, schedule, seeds, and all other hyperparameters fixed. Here we report its complete output on one dataset (Solar-Energy, $H=96$) together with two targeted probes on other datasets that changed the released configurations. We state the scope plainly: these are single-seed, single-horizon runs, and we therefore only interpret deltas well above the ≈ 0.005 MSE noise floor (Sect. 4.1); sub-floor deltas (e.g., the Mamba-vs.-MLP time encoder on solar) are reported as ties, and a full multi-seed ablation matrix across datasets and horizons is deferred to the camera-ready version.

Table 5 isolates each component on Solar-Energy. The architecture’s claims hold: removing either branch degrades accuracy (the dual-domain structure is not redundant), the learned gate beats plain averaging, the DLinear anchor is load-bearing, and the variate mixer is the single largest contributor (+0.022 MSE when removed) — consistent with 137 strongly coupled photovoltaic plants.

Normalization is a per-dataset decision. The same probe that revealed RevIN *hurting* solar (−0.016 MSE when removed) shows the opposite on traffic: removing instance normalization degrades traffic from 0.405 to 0.505 MSE at $H=96$. The two datasets look alike on the surface — both bounded, both strongly

periodic — but solar’s night zeros are *periodic structure* (window statistics then mostly encode the window’s night fraction, and de-normalization amplifies errors by an unstable standard deviation), whereas traffic windows carry genuine distribution shift (weekday/weekend composition, per-sensor levels) that instance normalization absorbs. We therefore ship normalization as a verified per-dataset switch and caution against transferring normalization conclusions across datasets.

Negative results that shaped the design. Deepening the variate mixer on traffic from two to four layers (19.7M $\rightarrow$ 37.9M parameters) produced no accuracy gain, so the released configuration keeps two layers; the residual gap to S-Mamba on traffic appears to require modeling the weekly cycle, which at hourly sampling (168 steps) exceeds the protocol look-back of 96 and is invisible to both branches. Similarly, enabling the variate mixer on ETT (7 channels, 8.4K training windows) tripled parameters and *caused* premature early stopping through immediate overfitting; the released ETT configurations disable it.

4.4 Analysis

Where each domain wins. The branch-only ablations (Table 5) quantify the two domains directly on solar: removing the time branch costs +0.012 MSE while removing the frequency branch costs +0.002, i.e., both contribute but the time view carries more — consistent with the protocol look-back hiding part of the periodic structure (at 10-minute resolution the daily cycle spans 144 samples $> L=96$; the same argument applies to weather’s daily and traffic’s weekly cycles). Beyond ablations, the fusion gate (Eq. 5) exposes per-sample, per-channel attribution of the two domains at no extra cost; a systematic study of gate behavior across datasets is left for the extended version.

Efficiency. DD-Mamba trains in under one minute per epoch on all nine datasets on a single consumer GPU. The bidirectional variate mixer scans C tokens rather than $B \times C$ folded time sequences, so high-channel datasets (862 variates) remain memory-light; the causal time scan is reserved for datasets whose channel count keeps it affordable.

5 Conclusion

DD-Mamba shows that a state-space forecaster gains from committing to *both* domains at once — provided each branch is anchored by the strongest linear model of its domain, deep components are zero-initialized corrections, and capacity (especially cross-variate mixing) is placed according to dataset characteristics rather than uniformly. Under the unified look-back-96 protocol, this recipe delivers best or matching average MSE on eight of nine standard benchmarks against the strongest published baselines, with parameter counts spanning 0.09M to 19.7M across datasets. Ablations verify every component and surface a transferable methodological caution: instance normalization can help or hurt

substantially depending on whether apparent nonstationarity is real shift or periodic structure. Limitations remain. Methodologically: main results are single-seed (we treat sub-noise-floor differences as ties, and multi-seed mean $\pm$ std reporting is deferred to the camera-ready), the complete ablation suite is reported on one dataset and horizon, and baseline numbers are quoted from [12] under an identical protocol rather than re-run in a unified framework. Architecturally: the fixed look-back of 96 hides cycles longer than the window (weather’s daily and traffic’s weekly periods), which we believe explains the remaining deficit on traffic; multi-resolution windows and adaptive spectral cutoffs are natural next steps.

Reproducibility. Code, configurations, and the ablation harness are provided as supplementary material and will be released in a public repository upon acceptance.

References

1. Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., Zhang, W.: Informer: Beyond efficient transformer for long sequence time-series forecasting. In: AAAI. pp. 11106–11115 (2021)
2. Wu, H., Xu, J., Wang, J., Long, M.: Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. In: NeurIPS. pp. 22419–22430 (2021)
3. Zhou, T., Ma, Z., Wen, Q., Wang, X., Sun, L., Jin, R.: FEDformer: Frequency enhanced decomposed transformer for long-term series forecasting. In: ICML. pp. 27268–27286 (2022)
4. Nie, Y., Nguyen, N.H., Sinthong, P., Kalagnanam, J.: A time series is worth 64 words: Long-term forecasting with transformers. In: ICLR (2023)
5. Liu, Y., Hu, T., Zhang, H., Wu, H., Wang, S., Ma, L., Long, M.: iTransformer: Inverted transformers are effective for time series forecasting. In: ICLR (2024)
6. Zeng, A., Chen, M., Zhang, L., Xu, Q.: Are transformers effective for time series forecasting? In: AAAI. pp. 11121–11128 (2023)
7. Li, Z., Qi, S., Li, Y., Xu, Z.: Revisiting long-term time series forecasting: An investigation on linear mapping. arXiv:2305.10721 (2023)
8. Kim, T., Kim, J., Tae, Y., Park, C., Choi, J.H., Choo, J.: Reversible instance normalization for accurate time-series forecasting against distribution shift. In: ICLR (2022)
9. Yi, K., Zhang, Q., Fan, W., Wang, S., Wang, P., He, H., An, N., Lian, D., Cao, L., Niu, Z.: Frequency-domain MLPs are more effective learners in time series forecasting. In: NeurIPS (2023)
10. Xu, Z., Zeng, A., Xu, Q.: FITS: Modeling time series with 10k parameters. In: ICLR (2024)
11. Gu, A., Dao, T.: Mamba: Linear-time sequence modeling with selective state spaces. arXiv:2312.00752 (2023)
12. Wang, Z., Kong, F., Feng, S., Wang, M., Zhao, H., Wang, D., Zhang, Y.: Is Mamba effective for time series forecasting? Neurocomputing (2025)
13. Ma, S., Kang, Y., Bai, P., Zhao, Y.-B.: FMamba: Mamba based on fast-attention for multivariate time-series forecasting. arXiv:2407.14814 (2024)

14. Wu, H., Hu, T., Liu, Y., Zhou, H., Wang, J., Long, M.: TimesNet: Temporal 2D-variation modeling for general time series analysis. In: ICLR (2023)